

CS 6650 Assignment 1

Overview

Through this assignment, you will learn how to build a client-server program which is the basis for any networked program including distributed systems. You will learn

- how to connect software components over the network using sockets,
- how to hide network details by building a simplified RPC-like abstractions using stubs,
- how to use threads and synchronization primitives to create multi-threaded client and server programs, and
- how to measure the basic performance metrics that can help you explore design trade offs.

Robot factory

You own a robot factory that produces robots on-demand. Your main customers are retailers who sell robots at their stores and the customers continuously place orders to your factory. You will create interfaces to receive orders and ship your robot, model the behavior of your customers, and manage your factory employees to create robots.

1 Create an order-and-delivery system (simple stubs)

For your factory to operate, customers should be able to place orders and receive robots. Similarly, your factory should be able to receive orders and ship robots. For this process, you as the factory owner and your customers do not need to know the very details of how orders and robots are delivered. You will implement a system that sends and receives orders and robots, but hide the low level communication details.

The client program models your customers and the server program is your factory. You should implement the low-level communication code using TCP sockets and hide them under high-level client and server wrappers. You should design `ClientStub` and `ServerStub` classes for the client and the server, respectively. Each client and server thread will later hold their own instances of these classes for communications.

The following describes the basic requirement and you can define more features and functions as needed.

1.1 Client stub

- `ClientStub.Init(std::string ip, int port)`: initializes the client stub and establishes a new TCP connection to the server.
- `ClientStub.Order(/* order details + any */)`: sends an order to the server and receives the robot (or robot information) from the server
 1. Once the function is called the order details should be marshalled into a byte stream.
 2. The byte stream is then sent through a socket connection to the server.
 3. Next, the function should wait for a server response.

4. Once the server responds back with robot information in a byte stream format, the byte stream should be unmarshalled to robot information and returned.

1.2 Server stub

- `ServerStub.Init(/* connected socket to a client*/):` initializes the server stub. It should directly take a socket or a class that includes a socket. In either case, the socket should be the one accepted from a listening socket and should be already connected to a client.
- `ServerStub.ReceiveOrder(/* any */):` receives order from the client
 1. The function waits for client order to arrive through the socket connection.
 2. Once the order is received through the socket in a byte stream format, the byte stream should be unmarshalled to an order and should be returned.

The order information will be used to assemble the robot (i.e., filling in robot information; see later sections). Once the robot assembly is finished, the robot information should be sent back to the customer.

- `ServerStub.ShipRobot(/* robot information + any */):` sends the robot information to the client who ordered it. It should marshal the robot information into a byte stream, and send the byte stream through the socket connection to the customer.

1.3 Orders and robot information

This section describes the information that must be included in the order and the robot information. Both order and robot information should be designed as objects using structs or classes.

Orders Each order must include,

```
int customer_id;    // customer id
int order_number;   // # of orders issued by this customer so far
int robot_type;     // either 0 - regular or 1 - special
```

Robot information Each robot information message must include,

```
int customer_id;    // copied from the order
int order_number;   // copied from the order
int robot_type;     // copied from the order
int engineer_id;    // id of the engineer who created the robot
int expert_id;      // id of the expert who added a special module
                   // -1 indicates that there is no special module
```

You may add extra information/variables if needed.

1.4 Hints

- It is recommended that you first build communication classes that encapsulates the socket-level code for clients and servers and then use these classes to build the client and server stub classes. You will be able to reuse the communication class for later assignments.

- Write a simple client and server program that simply sends and receives (order and robot info) messages to first test whether your communication class works.
- To test network communications while coding, you do not need two different computers. You can use the localhost IP address to test your client and server interactions within a single machine. Use "127.0.0.1" as the IP address and use port number of your choice (but do not use preassigned ports by other services such as 21, 22, 80, etc.).

2 Model your customers and factory

Now that you have the communication interface, you can interact with your customers. The next goal is to model your customer behaviors and to create your factory.

2.1 Customers (client program)

Your robot is very popular and each customer is only allowed to issue one order at a time one. Your robot sells out very quickly in your customers' retail stores and all of your customers place the next order immediately after they receive the robot from the factory.

Client program structure

You will model each customer as a thread and your client program will model multiple customers that concurrently issue orders to you factory. Your client program should take command line arguments that specify the ip address of the server, the port number of the server, the number of customers, how many orders each customer will place, and the robot type that the customers want. The client program should take these as arguments in the following format:

```
./client [ip addr] [port #] [# customers] [# orders] [robot type]
```

For example,

```
./client 123.456.789.123 12345 16 1000 0
```

means that the server ip is 123.456.789.123, the server port number is 12345, there will be 16 customers, and each customer will place 1000 orders, where the robot type is regular.

When the client program starts up,

1. The program should create the customer threads as many as the specified customer number.
2. Each customer thread should have a unique `customer_id`.
3. Either the main thread or each customer thread can instantiate connection to the server, but the socket connection should be made once per client stub and each customer should have its own client stub instance.
4. The customer thread should start issuing orders and receiving robot information as many times as the input argument using the client stub described in the previous section. The order should include the corresponding `customer_id`, `order_number`, and `robot_type`. For now, only use robot type 0 (regular type).
5. Once the thread completed all of its tasks, the connection to the server should be closed and the thread should terminate.

By varying the number of customers you should be able to control the amount of concurrent loads on the server.

Performance statistics

Each customer thread should measure how much time it took for each order (elapsed time from issuing an order to receiving the robot information) using `std::chrono::high_resolution_clock` in a **microsecond scale**. This time period is the latency for an order. Compute the mean latency, maximum latency, and minimum latency based on the measured time information for all orders. For example, if you have 100 customers who place 100 orders each, then you should have 10000 latency records to compute these numbers. In addition, measure the throughput of your factory (server program) from your client program (i.e., **orders/second**). Print these numbers at the end of each client program execution in the following tab separated format:

```
[avg latency]    [min latency]    [max latency]    [throughput]
```

2.2 Robot factory (server program)

Based on years of order history, you know that every new customer will place many orders. So you decided to hire and assign an engineer who produces robots for each customer. Currently, your robot factory only produces regular type robots.

Server program structure

The server program is your factory and threads in the server program are your engineers. Each engineer should have a unique id and once an engineer is assigned to a customer, the engineer should directly communicate with the customer.

The server program should take port number as a command line argument so that it can accept client connections through the port.

```
./server [port #]
```

For example, for a server to listen to new connection through port 12345, you should run,

```
./server 12345
```

The main thread of the server program should work as the following.

1. The server program creates sockets to accept new connections and waits for new connections from client programs in a loop.
2. Once a new connection is made, it creates a new engineer thread for the connection. The engineer thread should be given access to a server stub that uses the new connection to directly communicate with the customer thread.

Since your server waits for clients in a loop, you can type Ctrl+c in the terminal to terminate your server program.

Each engineer thread implements a loop that continuously processes orders until the client closes the connection. The thread uses the server stub to communicate with the client.

1. The thread waits for the client's order to arrive.
2. Once the order is received, it starts filling in the robot information for the order. It copies the information in the order to the robot information and adds its own `engineer_id` and `expert_id` (use value -1 for now). Then the robot information is sent back to the client.

If there the connection is closed by the client, the thread can escape the loop and terminate. A client closing the connection can be noticed by the return value of `socket recv()` function. When `recv` error occurs (including client connection close) `recv` returns -1. You can figure out more details about the `recv` errors by inspecting the `errno` (<https://man7.org/linux/man-pages/man2/recv.2.html>, <https://man7.org/linux/man-pages/man3/errno.3.html>).

3 Special robots

You were only selling regular robots and you want to add special robots to your production line. Special robots require an expert engineer who can add a special module to the regular robot. However, expert engineers are difficult to hire so you cannot assign them per customer.

Therefore, to create a special robot, engineers should create most of the robot and request the expert engineer to attach the special module. Because there are only a few expert engineers, requests should be queued and processed as the expert engineer becomes available.

Adding expert engineer workflow to robot factory

To add expert engineer workflow, the server program should take an additional command line argument, which is the number of expert engineers.

```
./server [port #] [# experts]
```

For example,

```
./server 12345 2
```

indicates that the server uses port number 12345 and has 2 expert engineers.

Expert engineers should be added to the server program as a thread pool. Recall that threads in a thread pool wait for tasks to arrive using guards (condition variables). You should use a single FIFO queue that is shared among all expert engineer threads and regular engineer threads. The regular engineers will enqueue requests to the queue and the expert engineers will dequeue the request, process it, and notify completion to the engineer who requested the task. The request sent to the expert engineer should include the robot information that the engineer was working on, and the response by the expert engineer should send back the robot information with the expert engineer's id filled in.

Once the regular engineer issued a request to the expert engineer, it should wait until an expert notifies completion of the request, and then send back the robot information back to the customer.

Each regular engineer thread should look at the customer order and send expert engineer requests only if the robot type is special. For regular robot types, the engineer should solely create the robot without expert engineer's involvement as described in Section 2.

Now, the main thread of the server program should be modified. The following indicates additional functions that you need to implement.

1. The server program creates the expert engineer request queue and the expert engineer thread pool. Expert engineer threads should be given unique ids that do not overlap with regular engineer ids.
2. The regular engineer threads should be given ways to access the expert request queue so that requests can be sent to the expert engineer threads.

The regular engineer thread should be modified so that it can communicate with the expert engineers, when needed. Necessary additions include,

1. Checking the robot type in the order and if it is a regular robot type, then follow the workflow in the previous section.
2. If the robot type is special, then send a request to the expert engineer threads by enqueueing the request to the expert engineer request queue. Next, wait for an expert engineer thread to respond back. Once the response is received, the regular engineer thread sends back the robot information with added expert id to the customer.

Each expert engineer thread in the thread pool should work as follows.

1. It waits for requests to arrive in the queue.
2. Once a request is detected and successfully received, the expert engineer thread works on the robot for at least 100 microseconds (**implement this using `std::this_thread::sleep_for()`**) and add expert id into the robot information. Then it respond back to the regular engineer thread that sent the request with the completed robot information.

Hints Note that the key components to support the expert engineer workflow are implementing:

1. A shared request queue.
2. An expert engineer thread pool.
3. Mechanisms for regular engineers to enqueue requests to the request queue safely.
4. Mechanisms for expert engineers to wake up and process the request only when there are pending requests in the queue.
5. Mechanisms for regular engineers to wait until any expert engineer processes the request.
6. Mechanisms for expert engineers to wake up the specific engineer who sent the request.
7. Mechanisms for expert engineers to notify their ids to the regular engineer who sent the request.

There can be many ways to implement the key components above. One way to implement a similar function to 1-4 was covered during the class. 5-7 requires your design skills which may involve using various synchronization primitives and shared variables. To make sure that the pair of regular engineer and the expert engineer who work on the same robot can exchange the wake up handles and ids, you can design the request to include rich information in addition to the robot information.

4 Playing with the implementation

Now that you have all the components, it's time to play with them and observe the performance trends that are reported from the client program. **Make sure your client and server programs are running on separate machines (see the next section for more details)**. Enter an appropriate number of orders so that your client program runs for at least 10 seconds. (e.g., place orders as many as 10,000.), but not more than 1 minute to save time and shared compute resources.

Experiment 1 Use the regular robot type and vary the number of customers: 1, 2, 4, 8, 16, 32, 64, 128, and 256. What are the measured mean latency, min latency, max latency, and throughput for each case?

Experiment 2 Use the special robot type and set the number of expert engineer to 1. Vary the number of customers: 1, 2, 4, 8, 16, 32, 64, 128, and 256. What are the measured mean latency, min latency, max latency, and throughput for each case?

Experiment 3 Use the special robot type and fix the number of expert engineer to 16. Vary the number of customers: 1, 2, 4, 8, 16, 32, 64, 128, and 256. What are the measured mean latency, min latency, max latency, and throughput for each case?

Experiment 4 Use the special robot type and set the number of expert engineer to be the same number as the customers. Vary the number of customers and expert engineers: 1, 2, 4, 8, 16, 32, 64, 128, and 256. What are the measured mean latency, min latency, max latency, and throughput for each case?

Plot the numbers For each experiment, plot two graphs (a total of 8 graphs): 1) a latency graph that includes the mean, min, and max latency numbers; and 2) a throughput graph. The y-axis should be the latency or throughput and the x-axis should be the number of customers. For each graph, write a couple of sentences that explain what is happening.

5 Running your client and server programs on Khoury VDI

Refer to the `vdi_setup.pdf` file to set up your VPN if you are off campus. Use your Khoury credentials to access the Khoury VDI machines. You may use VMWare Horizon Client as described in the document to use Linux machines with a GUI. If don't mind working without a GUI, you can directly ssh into the VDI machines. In either case, to run your client and server programs on separate machines, you will be using at least two terminals.

- If you are using VMWare Horizon Client, you can run, for example, the client program in your local terminal and use another terminal to run the server program on a different machine via ssh.
- If you are directly ssh'ing to the VDI machines you should ssh to two different VDI nodes.

You can access different VDI machines via ssh by specifying their addresses. VDI machine addresses range from `vdi-linux-030.ccs.neu.edu` to `vdi-linux-070.ccs.neu.edu`. **Refrain from using vdi-linux-030 machine and choose random numbers between 030 and 070 so that you and your peers are evenly distributed to different machines.**

Before running the client and server programs, make sure both programs are running on different machines. Type "hostname" in each terminal and it will return the machine name that your terminal is connected to. Type "hostname -i" to get the ip address of the machine that you will run the server program on. **Use port numbers between 10000 and 65535 for your server program:** these are freely usable ports between VDI machines.

Khoury VDI machines are shared resources, so be mindful of other users.

6 What to submit

1. Two-page report (11 pt font, single space) in a pdf format.

- Maximum one page summary of your software design, what works and what doesn't work, if any, and how to run your binary, if it is different from the specification above. You don't have to fill in the entire page. Concise and clear descriptions are the best.
 - 8 plots and their explanations in a single page as described in Section 4.
2. All source code files with a Makefile. The code should compile with a “make” command to create both “server” and “client” programs.

Place all files in a folder and create a zip file. Name the zip file “first name initials + last name.zip” (e.g., for Ji-Yong Shin create a file jshin.zip). Upload the file to the Canvas assignment section.

7 Due date

10/06/2020 (Tuesday), 11:59 pm. Refer to the syllabus on Canvas for the late policy.

8 Grading - max 15 points

1. The client and server stubs work (3 pt).
2. The client program works (3 pt).
3. The server program works for the regular robot type (3pt).
4. The server program works for the special robot type (3pt).
5. Two-page report (3pt).